

Understanding fine-tuning

Il y a 2 types de fine-tuning pour un modèle:

1. feature extraction

Dans cette méthode le modèle pré-entraîné a déjà appris des features donc on ne les change pas excepté la dernière couche qui est entraînée sur la tâche spécifique souhaitée

2. full-fine-tuning

Contrairement au feature extraction, le full fine-tuning entraîne le modèle dans son intégralité. Cette méthode exige plus de mémoire.

Il est possible d'aborder le fine-tuning de différentes manières:

- a. supervised fine-tuning:

- basic hyperparameter tuning: manually adjusting the model hyperparameters (learning rate, batch size...)

- transfer learning: a model pre-trained on a large, general dataset is used as a starting point. The model is then fine-tuned on the task-specific data, allowing it to adapt its pre-existing knowledge to the new task.

- multi-task learning: the model is fine-tuned on multiple related tasks simultaneously so that the model can develop a more robust and generalized understanding of the data. This works well when the tasks are closely related or when there is limited data for individual tasks.

- few-shot learning: In this method the model is given a few examples during inference to learn a new task. The goal is to guide the model by providing examples in the prompt. This method can be integrated into the reinforcement learning from human feedback (RLHF) if the small amount of data has human feedback

- task-specific fine-tuning: adapting the model parameters to the requirements of the targeted task thereby enhancing its performance to that particular domain.

- b. reinforcement learning from human feedback (RLHF)

- reward modeling: The model generates multiple outputs and human evaluators rank these outputs based on their quality.

- proximal policy optimization: Les mises à jour des poids sont minimisées pour éviter qu'elles soient trop éloignées pour éviter les changements abrupts et permettre une meilleure stabilité durant l'entraînement

- comparative ranking: similar to reward modeling but in this method the model learns from multiple rankings provided for each output (human evaluators). It focuses on the comparison between different outputs.

- preference learning: The model generates multiple outputs and the human indicates which output it prefers. Then the model adjusts its behavior to align with the human evaluator's preference.

→ parameter efficient fine-tuning: this method update a fraction of the model parameters by adding or modifying layers.